

State Exchange in Distributed Applications

ABSTRACT

This assignment is part of the TRAMP Real-time Application Mobility Platform (TRAMP) project with focus on distributed multimedia applications with real-time requirements. The goal of the assignment is to design and implement at least two components (producer and consumer) of a real-time application (such as a media player) and evaluate the provided distribution framework with respect to delay, throughput and other relevant metrics. Above is an overview of the provided framework. Your objective is to design and implement the producer and consumer application parts (the two gray components). The producer/consumer duo can run locally on one machine or be distributed over a network. In addition, multiple consumers can subscribe to the same produced data and receive identical copies.

Categories and Subject Descriptors

H.4 [Information Systems Applications]: Miscellaneous

General Terms

Design, Experimentation, Performance, Measurement

1. INTRODUCTION

In our assignment we were to use the Real-time Application Mobility Platform (TRAMP). This is a project developed by the Distributed Multimedia Systems (DMMS) group at the University of Oslo.

This is a assignment given to us in the course INF5090 - Advanced Topics in Distributed Systems which is a course given by the University of Oslo, as well as Lancaster University in England and University of Mannheim in Germany. The teachers are Thomas Plagemann and Vera Goebel with their teaching assistants Piotr Kaminski and Hans Vatne Hansen. These are also the people that developed the TRAMP framework together with some other developers all from the University of Oslo.

The TRAMP project was started in early 2010, the goal of the framework is to “make a migration system for real-time applications, such as Spotify and Skype, where users can take applications with them when they move.” [2]

Lets say you sit on your computer talking to a friend on skype. You need to go, but you got more to say to him/her. TRAMP will let you stream the conversation to your phone seamlessly and when you get to work, you can then transfer the conversation to your work computer.

2. RELATED WORK

2.1 Skype




Figure 1: Picture of the skype arcitecture [4]

As we've mentioned earlier in the report Skype is one of the applications the developers of the TRAMP framework reference to. Skype was the first peer-to-peer voice over

IP (VoIP) network, and requires minimal infrastructure in order to function. Skype came on market in late 2003 developed by Niklas Zennström and Janus Friis, the two main developer behind the flesharing system Kazaa. this was bought by eBay in late 2005 and again by Microsoft in May 2011. Since Skype is propriatory, the code is not available and this is therefore based on observations and partial reverse engineering of the protocol.

The skype architecture is designed around three entities: Supernodes, regular nodes and the login server (See fig. 1). Each regular node has a cache of the IP address and port-number to all reachable Supernodes. The Supernodes are a subset of the regular nodes. If a node has high uptime, good bandwidth and is not restricted by firewalls or Network Address translation (NAT) it can be chosen as a Supernode[4]. This will ofcourse put some extra stress on the nodes not behind NAT, but also make Skype possible to work for free due to the little infrastructure needed in order to make the network work. Another issue for these Supernodes is that they are used as third party for UDP hole punching to connect clients behind NAT. UDP hole punching is done by letting the regular nodes behind NAT connect to the third party (the Supernode) thereby opening ports that they can use for direct traffic between the two regular nodes. This is only open as long as there is communication traffic going. If there is a prolonged absence of traffic Skype sends "keep alive" packets instead of closing the connection and then having to use the Supernode as a third party again redo the connection.[3]

This is something that might be implemented in TRAMP to scale up the network.

2.2 Network-Integrated Multimedia Middleware (NMM)

<http://www.networkmultimedia.org>

3. SYSTEM DESIGN

The TRAMP platform is a platform designed to stream real-time data from a device to other devices connected to the session, mobility is the keyword. Lets take the example given in the introduction; You are on a skype call with some of your friends when you find out you need to go out of the house. Traditionally you then had to end the skypecall and call them up using your phone. The goal for TRAMP is to be able to seamlessly enter the conversation without your friend noticing you just changed device.

In TRAMP this is done by having a TRAMP daemon running on each device, with a consumer and a producer. The producer produces data, putting it into the shared memory, the TRAMP daemon then pushes the data to all other daemons connected to this session. When a TRAMP daemon gets data it puts it into the shared memory so the consumer can read it. All nodes has a TRAMP daemon, a producer and a consumer running. The shared memory of all nodes in the session are in synch.

4. THE CHOICE OF STREAMER

When we got this assignement we had to look at what we would implement in order to test the TRAMP platform. In

the sample code there was implemented a simple chat program, so we decided to implement something that streamed a lot of data. There are several ways to do this; we could make a file sharing system, an audio streaming application (like Skype) or a video streaming application. We decided on looking into FFMPEG.

FFMPEG is a complete cross-platform solution to record, convert and stream video and audio. It is a free software licensed under the LGPL or GPL (dependent on the configuration), and is the leading multimedia framework today.[1] FFMPEG provides tools for converting, streaming, playing and analysing multimedia, as well as a full developers library with the possibility to create allmoast anything you want. Big projects such as QStream, VLC, GStreamer and Google Chrome has used the framework.

Due to time constraints we decided we did not have enough time to implement this framework. Instead we implemented a streamer that copied over one big file by chunking the big file into small bits, sending them and putting them back together on the other side. This will in our eyes show the capabilities of TRAMP just as well as a more sophisticated streamer like the FFMPEG solution we had planned.

5. PROPOSED OPTIMIZATIONS

5.1 Shared Memory problem

5.2 Pushing more data than the recieving nodes can pull

5.3 UDP vs TCP or both

6. EVALUATION

7. CONCLUSIONS

8. REFERENCES

- [1] Ffmpeg-project homepage. April 2012.
- [2] Tramp-project homepage. April 2012.
- [3] D. K. Bryan Ford, Pryda Srisuresh. Peer-to-peer communication across network address translators. Technical report, Massachusetts Institute of Technology and Caymas Systems, Inc, 2005.
- [4] H. G. S. Salman A. Baset. An analysis of the skype peer-to-peer internet telephony protocol. Technical report, Department of Computer Science, Columbia University, New York, 2006.